

Consent Lens

Product Requirements Document — a comprehension layer for India's Account Aggregator consent moment

Rajeswari Rajagopalan

Version 0.2 · 19 August 2026 · pre-build

Author	Rajeswari Rajagopalan
Version	0.2 — pre-build
Date	19 August 2026
Status	Requirements agreed. Corpus not yet fetched. Every domain claim below is marked ⚠️ (secondary source, unverified) or ✓ (verified against primary source). At v0.2 almost everything is ⚠️ — that is the honest state of a document written before Step 1.
Scope	Business requirements. Architecture is out of scope except where a user-facing promise forces a structural constraint; those cases are marked and justified.

A note on this document's own discipline. Consent Lens exists to distinguish what a source actually says from what sounds right. A PRD for it that asserted unsourced domain facts would fail its own test. Section 14 records one instance where the first draft did exactly that.

Contents

1.	Summary
2.	Problem statement
3.	Users
4.	Why this is a problem for the user
5.	What the user actually cares about in that moment
6.	Requirements
7.	Implied requirements
8.	Non-negotiables
9.	The future user experience
10.	What will actually exist at build freeze
11.	How we would know it works
12.	Open decisions
13.	Risks
14.	Provenance note

1. Summary

India's Account Aggregator (AA) framework gives a citizen a legally sound, revocable, machine-readable way to share financial data. The pipe works. What does not work is the citizen's understanding of what they are agreeing to at the moment they agree to it.

Consent Lens is a retrieval-grounded assistant that answers a citizen's plain-language questions about an AA consent request, using only the published DEPA specification and the RBI Master Direction, and citing the clause every answer came from.

Its defining constraint is what it refuses to do. It explains and cites. It never advises, never adjudicates, never interprets the user's personal finances. That boundary is enforced in code, not in a prompt, and it is the product — not a safety wrapper around it.

2. Problem statement

2.1 The infrastructure is not the bottleneck

The AA framework crossed 100 million cumulative consents as of August 2024, with roughly 155 FIPs and 475 FIUs in the ecosystem at that point, after consent requests grew from 5.5M (FY23) to 63.75M (FY24). ⚠️ *Sahamati, Aug 2024 — two years stale; current figures to be refreshed before publication.* The rails are real and they carry volume. The precise number is not the argument; the order of magnitude is.

2.2 The artefact is written for machines, and shown to humans

An AA consent is a signed JSON object. Its exposure-determining fields are understood to include: ⚠️ *(field names and enums below are from FIU/AA integrator documentation, not yet from the primary ReBIT/DEPA spec — Step 1 confirms or corrects all of it)*

- **purpose** — a numeric code plus text label (101 wealth management, 102 spending/budgeting, 103 aggregated statement, 104 monitor accounts, 105 one-time access; the full ReBIT list may be longer) ⚠️
- **consentMode** — **VIEW / STORE / QUERY / STREAM** ⚠️
- **fetchType** — **ONETIME / PERIODIC** ⚠️
- **consentTypes** — **PROFILE / SUMMARY / TRANSACTIONS** ⚠️
- **fiTypes** — a longer enum including **DEPOSIT, MUTUAL_FUNDS, INSURANCE_POLICIES, EQUITIES** ⚠️
- **frequency** — pull cadence (**HOURLY / DAILY / MONTHLY / YEARLY** + count) ⚠️
- **dataLife** — retention period, including an **INF** (indefinite) option ⚠️
- **dataRange** — the historical window handed over ⚠️
- **consentStart / consentExpiry** — the life of the permission itself ⚠️

Four of these independently determine exposure, and they compose. **STORE + PERIODIC + DAILY + dataLife: INF** is a materially different act from a one-time read — and both render on a phone as a screen the user scrolls past in seconds.

The **dataLife / **consentExpiry** distinction is the sharpest illustration of the problem — and the single highest-value thing to get right in Step 1.** ⚠️ The working hypothesis is that revoking consent stops *future* data flow but does not by itself compel deletion of data already retained under **dataLife**; retention is governed separately.

If that holds, then a citizen who believes “I revoked, so they no longer have my data” holds an incorrect model of a system they legally consented to, and nothing on the screen corrects them.

This claim is contested and load-bearing. Post-revocation FIU deletion obligations may sit in the DPDP Act or FIU-side rules rather than the NBFC-AA Master Direction, which regulates Account Aggregators rather than FIU retention of received data. ⚠️ *If Step 1 shows the retention rule lives outside the corpus, that is itself a finding — it means the citizen’s most consequential question cannot be answered from the AA documents at all, and Consent Lens must abstain on it.* Either outcome is publishable. Assuming the answer is not.

2.3 Why comprehension matters more here than in other Indian DPI

UPI and Aadhaar have compulsion or overwhelming convenience behind them. **AA has neither. Nobody is required to use it.** That makes comprehension a direct adoption input rather than a downstream quality concern:

- A citizen who doesn’t understand the screen and abandons it is lost adoption.
- A citizen who doesn’t understand the screen and approves anyway produces a compliance artefact, not an informed decision — and the framework’s legitimacy rests on the claim that it is the latter.

Comprehension is an adoption bottleneck for infrastructure nobody is forced to use.

2.4 Where the gap is *not* the product’s to close

Some citizen questions are expected to be unanswerable because the *framework or its implementations* do not cover the case — not because retrieval failed. Candidates, all ⚠️ and all requiring a further distinction (see §14):

Case	Status ⚠️
Joint accounts	Reported as not supported for AA discovery across banks. Unresolved: spec-level exclusion or universal FIP implementation gap?
NRE / NRO accounts	Reported as not discoverable via AA
Non-sole-proprietor current accounts (partnership, company, trust)	Reported as excluded
fiTypes with no live production schema (bonds, debentures, G-secs, EPF, PPF)	Reported as proposed-not-live
FIP-AA fragmentation — a user’s bank and insurer may not be reachable through one AA	Reported

Telling these failure classes apart — *my retrieval failed* / *the corpus doesn’t contain it* / *the specification excludes it* / *the spec allows it but no FIP implements it* — is the primary analytical output of this build.

3. Users

	Who	Role
Primary	The citizen at the consent screen	Uses the product. No financial-services vocabulary, mid-task, deciding in seconds.
Secondary	The FIU or AA app that embeds it	Adopts it. Cares about drop-off, complaint volume, defensible consent quality.
Out of scope	Regulator, FIP, SRO	Beneficiaries of better consent quality. Not designed

Who	Role
	for.

The tension between primary and secondary is the commercial crux: **a comprehension layer that increases friction will not be embedded**. This constrains the product to be lateral and optional — a question the user may ask — never an added step in the flow.

4. Why this is a problem for the user

1. **The decision is consequential; the feedback is invisible.** A **PERIODIC** consent produces no further UI. Data flows monthly for a year and the user sees nothing. No wrong decision announces itself.
2. **The asymmetry is total.** The user authorises ongoing access by an entity they met ninety seconds ago, in vocabulary written for API integrators.
3. **Every other party on that screen benefits from a “yes.”** The FIU wants the data; the AA's revenue depends on completed data flows. ⚠️ (AA commercial models vary — commonly per successful fetch or FIU-borne fees rather than per consent granted; to be checked before this appears in the public note.) The user is the only participant with no advocate present.
4. **Rights that aren't understood aren't rights.** Revocation exists. A user who doesn't know it exists, or misunderstands what it reverses, holds a right they cannot exercise correctly.
5. **Language excludes exactly the intended beneficiaries.** The users AA is meant to bring into formal credit are the least served by English regulatory prose.
6. **There is nowhere to ask.** The lender's support line is an interested party. There is no neutral place to put the question, and asking a human carries social cost.

5. What the user actually cares about in that moment

Not what they'd say in a survey — what is true in the fifteen seconds. This section drives the prioritisation below, and each row names the requirement that implements it.

What they care about	What it demands	Implemented by
Finishing what they came to do. They came for a loan, not a data-rights education.	1–3 sentences. Never a wall of text. Never a required step.	FR-16
“How much of me is leaving?”	Scope expressible concretely — which accounts, which fields, how far back.	FR-4, FR-5
“Can I stop this?” Reversibility lowers the cost of being wrong; it is the highest-leverage single fact available.	Revocation and the retention/expiry distinction first-class, not buried.	FR-4, FR-7
“Am I being scammed?”	The product cannot answer this . It must decline the judgment and offer the scope-and-duration facts the user can judge on — without implying it has answered the safety question.	FR-6, FR-6a
Not being made to feel stupid.	No jargon returned in the answer. A machine is the right form factor precisely because asking it costs no dignity.	FR-16

What they care about	What it demands	Implemented by
Not being sold to.	See below.	FR-6

The insight that justifies the core constraint

Every party rendered on that consent screen benefits from the user pressing Approve. **An explainer that also recommended would be indistinguishable from the rest of the screen.** The refusal boundary is not primarily a liability control — it is the only thing that makes the explanation worth trusting. Neutrality is the entire value proposition, and neutrality that can be argued out of the system with a clever prompt is not neutrality.

Hence: **the boundary is enforced in code and demonstrable as a mechanism.**

6. Requirements

Priority key — **P0** must exist at build freeze, end of Sat 22 Aug 2026 · **P1** ships if the ~8–10 hour budget holds · **P2** explicitly deferred.

Budget gate: the brief allocates ~8–10 hours across Thu 20 – Sat 22 Aug. P1 items are cut in ascending ID order if P0 is not complete by the midpoint of Sat 22 Aug.

6.1 Functional

ID	Requirement	Pri
FR-1	Ingest the DEPA consent-artefact specification and the RBI Master Direction (NBFC-AA), preserving clause and section identifiers as metadata.	P0
FR-2	Chunk on clause/section boundaries rather than fixed token windows. <i>(Structural, not architectural: FR-5's promise that a user can open the source and find the claim is undeliverable if a citation points at half a clause.)</i>	P0
FR-3	Retrieve relevant clauses for a natural-language question.	P0
FR-4	Generate answers constrained to retrieved context. No answer may assert anything absent from the retrieved clauses.	P0
FR-5	Verifiable citation on every answer: source document, version and as-of date, clause/section identifier, quoted span. The user must be able to open the source and find it.	P0
FR-6	Refusal layer , enforced in code, classifying the question <i>before generation</i> , intercepting (a) advice — “should I approve”, (b) adjudication — “is this lender safe/legitimate”, (c) personal-finance interpretation. The evaluative question is never passed to the generator as a question to answer.	P0
FR-6a	Reformulation on refusal. Where an intercepted question contains an answerable factual sub-question, the system declines the judgment and separately answers the factual part through the normal grounded, cited path — without implying the judgment was answered. <i>(This is what makes refusal useful rather than a dead end; it is also the mechanism that satisfies §5's “Am I being scammed?” row without violating FR-6.)</i>	P0
FR-7	Abstention — distinct from refusal, with distinct wording. When retrieval returns nothing adequate, state that, and state which case applies: not in this corpus / excluded by the specification / spec-permitted but not implemented by FIPs / unresolved.	P0

ID	Requirement	Pri
FR-8	Failure log from the first query, recording question, answer, defect, and classification: retrieval / generation / corpus gap / framework gap . Target 15–20 entries. Written from entry one to be read by someone else.	P0
FR-9	Interface sufficient to demonstrate the above. CLI is acceptable.	P0
FR-10	Self-description — the system can state what it does and does not do, so a refusal reads as designed behaviour rather than malfunction.	P0
FR-16	Answer form constraints, enforced not merely prompted: hard length cap (≈3 sentences / ~60 words), and no unglossed jargon — any spec term appearing in an answer must be accompanied by its plain-language meaning.	P0
FR-17	Framework-gap register — a curated, sourced list of cases known to be outside the framework or unimplemented, consulted when retrieval is empty. <i>(Without it, FR-7's second branch is undeliverable: an empty retrieval is indistinguishable from a framework gap.)</i> Seed from §2.4; every entry carries its evidence and its ⚠️/✓ status.	P0
FR-11	Answer output in one Indian language. Retrieval on the English source; the <i>answer</i> is translated; the citation stays in the source language and is marked as authoritative.	P1
FR-12	Golden set of ~20–25 questions with expected clause, including questions whose correct outcome is refusal and questions whose correct outcome is abstention.	P1
FR-13	Adversarial prompt set attempting to argue the refusal layer out of its boundary; results logged.	P1
FR-14	Public GitHub repository: code, this PRD, the unedited failure log, the product note, README that runs from a clean clone.	P1
FR-15	Log translation degradation as a finding — what meaning is lost when a legal clause becomes plain vernacular.	P1
FR-18	Two-page product note, written by the author. Structure per brief §7. Delivered as PDF or DOCX. Assistant support is limited to drafting help, structuring and critique — the note is not to be ghost-written; it must be in the author's voice, because it will be defended in person.	P1
FR-19	One synthetic consent artefact to ask questions against, so demonstrations are concrete rather than abstract. <i>(Dependent on D-3; §9's scenes assume it.)</i>	P1

6.2 Non-functional

ID	Requirement	Pri
NFR-1	Every layer explainable by the author to a hands-on technologist. No framework that hides the retrieval step. Outranks convenience and outranks quality gains from opaque tooling.	P0
NFR-2	No real user financial data, in any environment, at any point. The only artefact the system parses is synthetic and committed to the repository; there is no ingestion path from a live AA, FIP or FIU. <i>(Note the honest limit: the parser cannot tell synthetic from real. The guarantee is operational — no live integration exists — not cryptographic.)</i>	P0
NFR-3	Corpus read-only and versioned. Every chunk carries source version and ingestion date.	P0

ID	Requirement	Pri
NFR-5	Logged queries contain no third-party PII; the log is a public artefact.	P0
NFR-4	Reproducible from a clean clone with a documented API key.	P1
NFR-6	Answer latency plausible at a consent screen. Measured and reported, not engineered for.	P2

6.3 Explicitly not building

Each is a deliberate decision; the reasoning is part of the deliverable.

Not building	Why
Live AA API integration / real financial data	Integration risk and regulatory surface; proves nothing about comprehension. A synthetic artefact demonstrates the same thing.
A redesigned consent flow	A larger, different product. Redesigning the screen without evidence of what citizens misunderstand is guessing.
Any UI worth looking at	Polish is the least informative use of the remaining hours.
DPDP Act / consent-manager overlap	Important, unresolved, and would consume the entire budget. Note where it touches §2.2.
Voice	Probably the correct eventual answer for this user base. Not decidable in three days.
Evaluation framework beyond the golden set	A hand-built set I can defend beats a framework I cannot.
A second or third corpus (UPI advisories, Aadhaar)	The thesis is one comprehension layer, many corpora. Proving it twice proves nothing new.

7. Implied requirements

Not stated in the brief; they fall out of the above and are easy to miss until they bite.

1. **Refusal ≠ abstention ≠ error.** Three states, three user-visible messages. Collapsing them makes a careful system look evasive or broken.
2. **Classification precedes generation.** Otherwise the generator receives an evaluative question alongside tempting context.
3. **Citations need a stable anchor.** “Verifiable” constrains ingestion — clause number, page or URL fragment must survive the pipeline. A UX promise that reaches back into parsing.
4. **Citations need an as-of date.** The Master Direction is amended. An undated citation is unverifiable later, and the argument for RAG over fine-tuning collapses if the citation cannot be dated.
5. **Translated answers must not read as translated law.** Vernacular output is a comprehension paraphrase; the authoritative text stays in the source language and is visibly marked so.
6. **“No answer” must be a passing outcome in the golden set** — else the set silently optimises toward always answering, the exact failure mode being guarded against.

7. **The refusal layer needs adversarial tests.** “Enforced in code” is an unproven claim until something has tried to break it.
8. **The framework-gap response needs a designed shape** (FR-17). It is the product’s best moment and should not be improvised at runtime.
9. **The corpus needs a stated boundary.** The system must know what lies outside DEPA + RBI MD, or it will answer FIU-specific questions from parametric knowledge.
10. **A gap needs a cause, not just a label.** “Framework gap” is too coarse: spec exclusion, absent FIP implementation, and rule-lives-in-another-instrument are three different findings with three different implications.

8. Non-negotiables

Drop any of these and the thing being demonstrated no longer exists.

1. **It explains. It never advises, and it never adjudicates.** Enforced in code. Demonstrable as a mechanism.
2. **No answer without a resolvable citation.** Abstain instead. An uncheckable citation is worse than none. *(Applies to claims about the corpus. Statements about the system’s own boundaries — “I don’t assess lenders” — are not corpus claims and are exempt; they must never be dressed as citations.)*
3. **No real user financial data, in any environment, at any point.**
4. **The failure log ships unedited.**
5. **The author can explain every layer.** Any component that cannot be explained is removed regardless of what it adds.
6. **One corpus.** The scope line does not move.
7. **The date is fixed; features are the variable.** If the budget is threatened, features are cut — never the failure log, never the note.
8. **No unsourced domain assertion reaches a public artefact.** ⚠ or ✓ on everything. This document is the first test of that rule (§14).

9. The future user experience

Consent Lens as a product, once it exists. §10 states plainly what will and will not be true on 24 August.

*Citations below use **[clause id]** placeholders. Real identifiers do not exist until the corpus is fetched, and inventing plausible ones in a document about verifiable citation would be self-refuting. The shape is the specification; the IDs arrive in Step 1.*

Lakshmi, 34, runs a two-person tailoring unit in Coimbatore. She is applying for a ₹2 lakh working-capital loan in a lender’s app and has reached the consent screen. Her Tamil is fluent, her English functional, her financial-regulatory vocabulary nonexistent. She wants the loan; a form stands between her and it.

Scene 1 — the fifteen seconds. The screen shows purpose 104, **fetchType: PERIODIC, dataLife: 6 MONTH**, granted 14 August 2026. Beside it, small: *Not sure what this means? Ask.* She taps.

“What am I actually giving them?”

Your savings account transactions from the last 12 months, and an updated copy every month while this consent is active. Not your PAN, not your address, and no account you haven’t linked. — DEPA consent artefact spec, \$Data Request Details · [version] · retrieved [date]

“For how long?”

They can collect new data until 14 February 2027. Separately, they may keep what they’ve already received until 14 August 2027. Those are two different clocks. — [source] [clause id] · retrieved [date]

Two questions, eleven seconds, no jargon returned to her. She approves, now holding a roughly correct model of what she did — which she did not sixty seconds ago.

Scene 2 — the refusal, which is the point.

"Is this lender safe? Should I approve?"

I can't answer that. I only explain what the consent document says — I don't assess lenders and I don't recommend decisions.

What I can tell you is what this request covers: 12 months of savings transactions, refreshed monthly until 14 February 2027, retained up to 14 August 2027. Whether that's acceptable is your call.

The evaluative question was intercepted and never reached the generator; the factual sub-question was re-issued through the normal cited path (FR-6a). It has refused, said why, handed her the material facts — and has not become another voice on the screen arguing for yes.

Scene 3 — four months later. Her loan is disbursed. She wonders whether data is still flowing.

"I want to stop this. Will they delete what they already have?"

You can revoke from your AA app; new data collection stops immediately.

On deletion: this consent's retention period runs to 14 August 2027. Whether revoking obliges them to delete sooner is not something the AA documents settle — that sits in separate data-protection rules I don't cover. Your AA app or the lender's grievance officer can tell you. — [source] [clause id] · retrieved [date]

This is the answer she could not get anywhere else. Note what it does *not* do: having found the corpus silent on the deletion question, it says so and points somewhere real, rather than producing a confident sentence. **If Step 1 shows the corpus does settle this, the answer gets shorter and gains a citation. Either way the shape is honest.**

Scene 4 — the honest gap.

"What about the account I hold jointly with my husband?"

Joint accounts can't currently be linked or shared through AA — as far as I can tell, no bank supports it. I should be straight with you about what I know: I can see that it doesn't work in practice. Whether that's because the rules exclude joint accounts or because banks simply haven't built it, I can't tell you from these documents.

No hallucinated clause. No confident nonsense. The system distinguishes *I don't know* from *the infrastructure doesn't do this* — **and distinguishes both from *I know it fails but not why*, which is the true state of this particular question.**

What changed for Lakshmi: she made the same decision she'd have made anyway, but knowingly; she knows how to reverse it; and she carries one true thing about the system's limits to the next screen. **What changed for the ecosystem:** consent quality became observable, and the gap between "legally valid consent" and "informed consent" narrowed by the width of four questions.

10. What will actually exist at build freeze

Stated plainly, because the gap between §9 and this is itself the prioritisation evidence.

Committed (P0, Sat 22 Aug): a command-line assistant answering plain-language questions over the real DEPA and RBI corpus with verifiable clause citations; a refusal layer enforced in code, with reformulation; abstention distinguished from refusal and gaps attributed to a cause; a framework-gap register; answer-form constraints; an unedited failure log of 15–20 classified entries.

Intended but budget-dependent (P1, cut in ID order if P0 slips): the second language, the golden set, adversarial results, the public repository, translation-degradation findings, the product note, the synthetic artefact. *The note and the repository are the two P1 items whose loss would most damage the interview; if the budget binds, cut FR-12, FR-13 and FR-15 first.*

Will not exist: any AA integration, any real user data, any phone interface, any measurement of real comprehension outcomes.

The claim is not “I built a product.” It is: *I found a real bottleneck in public infrastructure I know well, built the smallest honest probe of it, drew a hard boundary and enforced it in code, and learned more from the failures than from the successes.*

11. How we would know it works

Deliberately not “accuracy.” Accuracy over a hand-built set of 25 questions means whatever the set was designed to mean.

In this build:

- **Uncited-assertion rate** — of answers given, how many contain a claim about the corpus not supported by the attached citation. Target zero. (*Measured by manual audit of every golden-set answer. Not “citation present” — citation presence is guaranteed by construction and measures nothing; the question is whether the citation supports the sentence.*)
- **Refusal recall** — of adversarial advice/adjudication prompts, how many were caught.
- **Over-refusal count** — how often it refused something it should have answered. *This number will look bad, and it should be reported. Over-refusal is the correct failure direction.*
- **Abstention honesty** — of questions with no corpus answer, how many were abstained on rather than answered.
- **Gap attribution accuracy** — of known-excluded cases, how many were correctly attributed to the right cause (spec exclusion / no FIP implementation / rule elsewhere / unknown).

In production — what would actually matter, none of it measurable here:

- **Comprehension outcome** — can a user, one minute after consenting, correctly state what they shared, for how long, and how to stop it? The only metric that measures the actual problem.
- **Revocation rate.** ⚠ *Believed low; not sourced.* If low revocation reflects ignorance rather than satisfaction, a rise is a success signal — but that inference needs a baseline this build cannot produce.
- **Drop-off at the consent screen, split by Consent Lens engagement.** Both directions are informative: fewer confused abandonments is good; *more informed refusals* is also good. A product that optimises only the first is back to arguing for yes.
- **Repeat consent quality** — do users who used it once configure narrower consents next time?

12. Open decisions

#	Decision	Options	Recommendation
D-1	Which Indian language	Hindi / Tamil	Tamil. Weaker model coverage and richer morphology make degradation more visible — a better finding — and the author can verify output directly. Hindi is the safer demo. A choice between a better result and a better story;

#	Decision	Options	Recommendation
			the brief says judgment over polish.
D-2	Publication timing	Repo + post Thu 20 / repo Thu, post Sun 23	Split — see §13, risk 1.
D-3	Synthetic consent artefact (FR-19)	Yes / no	Yes, one. ~30 minutes; makes every demo question concrete. §9 assumes it.
D-4	Corpus: full RBI MD or consent-relevant clauses only	Full / subset	Decide after inspection. A subset is defensible if the boundary is documented.
D-5	If Step 1 shows post-revocation deletion is governed outside the corpus	Extend corpus / abstain and log	Abstain and log. Extending the corpus breaks non-negotiable 6. The abstention is a better finding than the answer.

13. Risks

#	Risk	Severity	Mitigation
1	Publishing before the refusal layer and failure log exist. The post's entire thesis is "it refuses, and here is how it failed." Posting Thu 20 Aug means publishing a demo — precisely the novelty-over-judgment framing to be avoided.	High	Make the repository public Thu 20 Aug with this PRD and a README carrying the thesis; that alone is referenceable. Post to LinkedIn Sun 23 Aug evening — after the log exists, and ~12 hours before the interview rather than four days stale in the feed.
2	This document is itself a public artefact containing unverified domain claims , read by an AA expert possibly before the interview.	High	Nothing unmarked. Refresh §2.1 figures and resolve every ⚠ in §2.2 and §2.4 before the repo goes public Thu 20 Aug; downgrade anything still unresolved to an explicit open question. §14 records that this risk already materialised once.
3	RBI Master Direction is a long PDF; clause-boundary chunking harder than expected	Medium	Time-box ingestion. If clause extraction fails on the MD, ship with DEPA fully structured and the MD page-anchored, and log it.
4	Refusal layer over-refuses; demo looks broken	Medium	Report over-refusal honestly. It is the correct failure direction and saying so is stronger than hiding it.
5	Repo + LinkedIn post are author-added scope, absent from the brief's 8–10 hour budget. 23 Aug is reserved for interview mocks.	Medium	Treat as ~90 minutes total, non-build, outside the budget. If it threatens P0, the post goes and the repo stays.
6	Scope creep into a real UI	Medium	§6.3 is the contract.
7	Corpus fetch blocked or spec unavailable at	Low	Snapshot everything on day one; commit the

#	Risk	Severity	Mitigation
	published URLs		snapshot with retrieval dates.

14. Provenance note

The first draft of this PRD asserted, without a marker and as settled fact, that joint accounts are excluded from the AA *framework* — and cited an invented clause number for a retention rule that may not sit in the AA Master Direction at all.

Both errors are exactly the failure mode Consent Lens is built to detect: a plausible sentence, confidently delivered, traceable to no source, conflating *the rules exclude this* with *nobody has implemented this*.

They are recorded here rather than quietly corrected, and they become **failure log entries 0a and 0b** — logged before the system ran its first query. If a careful human drafting slowly makes this error, an LLM retrieving quickly will make it too, and the distinction in FR-7 and FR-17 is not academic.

Appendix A — Corpus

Source	Location
DEPA consent artefact specification	https://depa.world/learn/consent-artefact
DEPA docs repository	https://github.com/depa-world/docs
RBI Master Direction — NBFC Account Aggregator	rbi.org.in , current consolidated version
ReBIT AA technical specifications	AA API / consent schema specs
Sahamati (RBI-recognised SRO) + FAQ	https://sahamati.org.in/ · https://sahamati.org.in/faq/

Appendix B — Framework-gap register, seed

Seeds FR-17. Every entry needs evidence, a cause attribution, and a status before it may be used in an answer.

Case	Cause attribution	Status
Joint accounts	Unresolved — spec exclusion vs universal FIP non-implementation	⚠ secondary source
NRE / NRO accounts	Unresolved	⚠ secondary source
Non-sole-proprietor current accounts	Unresolved	⚠ secondary source
fiTypes without live schema (bonds, debentures, G-secs, EPF, PPF)	Likely implementation, not spec	⚠ secondary source
FIP-AA fragmentation	Implementation / commercial	⚠ secondary source
Post-revocation deletion obligation	Possibly governed outside the AA corpus	⚠ unresolved — see D-5

Markers: ⚠ = asserted from secondary sources, unverified against the primary document · ✓ = verified against primary source with clause and retrieval date. No ⚠ item may appear in the product note or in a system answer without either being resolved to ✓ or being stated to the user as unresolved.